

Lecture 3 Supplement 1

01211373 Machine Learning and Programming for Industry

Dr. Varodom Toochinda
EMME, Kasetsart University

Lecture 3 supplement 1



Linear Regression

Notebook Supplement : `mlpi26_lect02_nbsup1.ipynb`

AI's explanation on the term "regression"

- "The term "regression" in statistics comes from Sir Francis Galton's observations about heredity, specifically the tendency of offspring to regress towards the mean of a population. In the late 1800s, Galton studied the relationship between parents' and children's heights and noticed that exceptionally tall or short parents tended to have children whose heights were closer to the average height of the population. This "regression towards the mean" is the origin of the term, which later evolved into the broader statistical concept of regression analysis."

Reference videos for linear regression

- Stanford CS229 ML course (Andrew NG 2018), Lecture 2 : Linear Regression and Gradient Descent
https://youtu.be/4b4MUYve_U8?si=qa3iBTO2R-abQViu
- MIT: Machine Learning 6.036, Lecture 5: Regression (Fall 2020)
<https://youtu.be/OpvrwhWEUmY?si=bS8VZYsrZMqvnuii>

Simple Linear Regression

(see lecture3.ipynb for more)

- In general, find $h_{\theta}(x)$ that maps x in $\mathbb{R}^d$ to y in $\mathbb{R}$ (d is dimension of feature vector)
- $h_{\theta}(x)$ is called a hypothesis or model
- In linear regression literature, $\theta_i, i = 1, 2, \dots$ is often used in place of w_i and θ_0 for b .
- First consider simple case $h_{\theta}(x) = \theta_0 + \theta_1 x_1$
- Define intercept term $x_0 = 1$ so that
 - $h_{\theta}(x) = \sum_{i=0}^d \theta_i x_i = \theta^T x$
- Define cost function

$$J(\theta) = \frac{1}{2} \sum_{i=1}^n (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

LMS algorithm

Objective : Choose θ to minimize $J(\theta)$

Gradient descent

- start with some initial guess for θ
- Update θ to minimize $J(\theta)$

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta), \quad j = 0, \dots, d$$

Here α is called *learning rate*.

It is shown in [2] that

$$\frac{\partial}{\partial \theta_j} J(\theta) = (h_{\theta}(x) - y)x_j$$

Hence, the LMS update rule for single training example becomes

$$\theta_j := \theta_j + \alpha (y^{(i)} - h_{\theta}(x^{(i)})) x_j^{(i)}$$

a.k.a **Widrow-Hoff** learning rule.

LMS Algorithm

For the whole data set, the algorithm becomes


Repeat until convergence {

$$\theta_j := \theta_j + \alpha \sum_{i=1}^n (y^{(i)} - h_{\theta}(x^{(i)})) x_j^{(i)}, \quad j = 0, \dots, d$$

}

Or, compactly with θ in vector form

$$\theta := \theta + \alpha \sum_{i=1}^n (y^{(i)} - h_{\theta}(x^{(i)})) x^{(i)}$$



Result from
linear regression

Normal equation

- The value of θ that minimizes $J(\theta)$ is given in closed form by

$$\theta = (X^T X)^{-1} X^T y$$



The background of the slide features several thin, curved lines in a light gray color, some solid and some dashed, creating a sense of motion or a stylized globe. A large, solid red rectangle is positioned on the left side of the slide, containing white text. The rectangle has a small triangular pointer at its bottom center.

Example from
chapter 9 of [1]

- See `lecture3.ipynb`

The background of the slide features a series of thin, curved lines in a light gray color, creating a sense of motion or a stylized globe. These lines are more prominent on the left side and fade towards the right.

Optional reading for linear regression from [2]

- Probabilistic interpretation of linear regression
- Locally weighted linear regression

1. Fit θ to minimize $\sum_i w^{(i)} (y^{(i)} - \theta^T x^{(i)})^2$
2. Output $\theta^T x$.

Locally weighted
linear regression

Standard
linear regression

1. Fit θ to minimize $\sum_i (y^{(i)} - \theta^T x^{(i)})^2$
2. Output $\theta^T x$.

Locally weighted linear regression
(non-parametric algorithm)

$$w^{(i)} = \exp\left(-\frac{(x^{(i)} - x)^2}{2\tau^2}\right)$$

References




1. S. Raschka, Y. Liu and V. Mirjalili. Machine Learning with PyTorch and Scikit-Learn. Packt Publishing. 2022.

<https://sebastianraschka.com/blog/2022/ml-pytorch-book.html>



2. A. Ng. CS229 Lecture Notes. Stanford University. 2023.

https://cs229.stanford.edu/main_notes.pdf



3. T. Broderick. Introduction to Machine Learning course (6.036). Massachusetts Institute of Technology. 2020.

<https://tamarabroderick.com/ml.html>